

1 Data Structures :

- **Series:** 1D labeled array, like a column
 - `s = pd.Series([1, 2, 3])`
- **DataFrame:** 2D labeled data structure, like a table with rows and columns.
 - `df = pd.DataFrame({'A': [1, 2], 'B': [3, 4]})`

Reading & Writing Data :

- `df = pd.read_csv('file.csv')` # Read CSV
- `df.to_csv('output.csv', index=False)` # Write CSV
- `df = pd.read_excel('file.xlsx')` # Excel
- `df.to_excel('output.xlsx', index=False)`

Viewing and Inspecting Data :

- `df.head()` # First 5 rows
- `df.tail()` # Last 5 rows
- `df.info()` # Info on columns, non-null count, dtypes
- `df.describe()` # Summary stats for numeric columns
- `df.shape` # Rows and columns
- `df.columns` # List column names
- `df.dtypes` # Data types of columns

Data Selection :

- `[] [] ; loc ; iloc`
- Select * from data
 - `movies[:,:]`
 - `movies.loc[:,:]`
 - `Movies.iloc[:,:]`
- select first 5 rows from the data using all methods
 - `loc` - Label (row/column names), Yes for slicing (end included)
 - `iloc` - Integer position (row/column index), No for slicing (end excluded)
 - `movies[0:5][:]`
 - `movies.loc[0:4,:]`
 - `movies.iloc[0:5,:]`
- select first 5 rows, 2 columns in the dataset
 - `movies.iloc[0:5,0:2]`
- select first 10 rows, and type, title, director columns
 - `Movies[0:10][["type", "title", "director"]]`
 - `movies.loc[0:9, ["type", "title", "director"]]`

Filtering Data :

- select * where type = movie and country = India
 - `movies.loc[(movies["type"]=="Movie") & (movies["country"] == "India"),:]`
- Select distinct country
 - `movies.loc[:, "country"].unique()`

- Select distinct country,type combinations
 - `movies.loc[:,["country","type"]].drop_duplicates()`

Handling Missing Values :

- Get the number of missing values in each column
 - `movies.isnull().sum()`
 - `movies.isnull()` gives true false values in all the cells. It has the same shape as of the dataframe
 - `movies.isnull().sum()` counts the true values
- Drop all rows containing missing values
 - `movies.dropna()`
- Replace all missing (NaN) values in the DataFrame df with 0
 - `movies.fillna(0)`

Iterating :

- `iterrows()` is a Pandas DataFrame method that lets you iterate over the rows of a DataFrame one by one.
- It returns each row as a (index, Series) pair.
 - Below prints each row one by one. Each row is a pandas series
 - `for index, row in movies.iterrows():`

```
print(index)
print(row)
print(type(row))
```
- `dict.items()` - Returns a view of (**key, value**) pairs in the dictionary.
 - `student = {"name": "Alice", "age": 25}`
`for key, value in student.items():`
`print(key, "→", value)`
 output
 name → Alice
 age → 25

Aggregations :

- select country, count(showid) from movies where type = movie
 - `df =movies.loc[movies["type"]=="Movie",:]`
`df.groupby("country")["show_id"].count()`
- select country, release year count(showid) as show_count, count(director) as director_count, countd(director) as director_distinct_count from movies where type = movie
 - `result = df.groupby(["country","release_year"]).agg({"show_id" : ["count"], "director" : ["count","nunique"]})`
`result.columns = ["show_count","director_count","director_distinct_count"]`
`result.reset_index()`

Merging, and Joins :

- `select * from movies where country = india union select * from movies where country = 'United States'`
 - `df1 = movies.loc[movies["country"]=="India",:]`
`df2 = movies.loc[movies["country"]=="United States",:]`
`df3 = pd.concat([df1,df2])`
`df3`
- create two dataframes, and join them (inner, left, right, outer)
 - `movies_sample = pd.DataFrame({`
 `'movie_id': [1, 2, 3, 4],`
 `'title': ['Inception', 'Titanic', 'Avatar', 'Interstellar'],`
 `'director_id': [101, 102, 102, 103]`
 `})`
`directors = pd.DataFrame({`
 `'director_id': [101, 102, 104],`
 `'name': ['Christopher Nolan', 'James Cameron', 'Steven Spielberg']`
 `})`
`inner = pd.merge(movies_sample, directors, on='director_id', how='inner')`
`inner`

Sorting :

- `select year, country, type, show_id from movies order by year desc, country asc`
 - `movies.loc[:,["release_year","country","type","show_id"]].sort_values(by = ["release_year","country"], ascending = [False,True])`
- Rank, Dense Rank, Row Number

Apply and Lambda

- The `.apply()` method applies a function to each group (if used with `groupby`), or to each row/column (if used on a `DataFrame`), or element (if used on a `Series`).
- Understanding what is getting passed to `apply` is important.
- `df.apply(lambda x : some logic) -->` logic gets applied on each column of `df`
 - Converts varchar columns to uppercase
 - `df = pd.DataFrame({`
 `'name': ['Alice', 'Bob', 'Charlie'],`
 `'id' : ['a','b','c'],`
 `'math': [80, 90, 70],`
 `'science': [85, 88, 95]`
 `})`

`df.select_dtypes(include='object').apply(lambda x: x.str.upper())`
- Same code can be written this way as well to see what is happening, and what is passed to `x`
 - `def uppercase(x):`

```

print(x)
print(type(x))
print(x.str.upper())
print("end")

```

```

df.select_dtypes(include='object').apply(uppercase)

```

- Both `.map()` and `.apply()` can apply a function to each value in a Series. But `map()` is slightly faster with series
 - `.apply()` does not accept a dictionary when used on a pandas Series. Only `.map()` supports direct dictionary mapping.
 - `titanic["sex"].map({'male': 0, 'female': 1})`
 - `.map()` is designed to map values directly using a dictionary, Series, or function.
 - `.apply()` expects a function, not a dictionary.
- `df["name"].apply(lambda x: x.upper())`
 - Here each value of the name column gets passed one by one
 - We can write a similar function like above to see that.
- Assign 1 if a , 2 if b, else 0 on id column
 - ```
def assign_flag(x):
 print(x)
 print(type(x))
 if x == 'a':
 print("end")
 return 1
 elif x == 'b':
 print("end")
 return 2
 else:
 print("end")
 return 0
df["id"].apply(assign_flag)
```
- `df.apply(lambda x : some logic, axis = 1) ---->>` logic gets applied on each row of df. Here each row gets passed one by one.
  - Try one example
  - We can write functions similar to above
- Apply with groupby
  - ```
df = pd.DataFrame({
    'Department': ['HR', 'HR', 'IT', 'IT'],
    'Employee': ['Alice', 'Bob', 'Charlie', 'David'],
    'Salary': [50000, 55000, 70000, 72000]
})
```
 - `df.groupby('Department')['Salary'].apply(lambda x: x.mean())`
 - ```
def group_dep(x) :
 print(x)
```

```
print(type(x))
print("end")
return x.mean()
mean = df.groupby('Department')['Salary'].apply(group_dep)
print(mean)
```

- `groupby('Department')` splits the DataFrame into two groups: HR and IT.
- `.apply(lambda x: x.mean())` computes the mean salary for each group.
- `.applymap()` is used to apply a function elementwise to every cell in a DataFrame only.
  - It does not work with Series or GroupBy objects.